

Homework 2

February 18, 2021

1 Codes for Q1

```
[ ]: from __future__ import absolute_import
from __future__ import print_function
from future.standard_library import install_aliases
install_aliases()
import numpy as np
from scipy.special import logsumexp
import os
import gzip
import struct
import array
import matplotlib.pyplot as plt
import matplotlib.image
from urllib.request import urlretrieve

def download(url, filename):
    if not os.path.exists('data'):
        os.makedirs('data')
    out_file = os.path.join('data', filename)
    if not os.path.isfile(out_file):
        urlretrieve(url, out_file)

def mnist():
    base_url = 'http://yann.lecun.com/exdb/mnist/'

    def parse_labels(filename):
        with gzip.open(filename, 'rb') as fh:
            magic, num_data = struct.unpack(">II", fh.read(8))
            return np.array(array.array("B", fh.read()), dtype=np.uint8)

    def parse_images(filename):
        with gzip.open(filename, 'rb') as fh:
            magic, num_data, rows, cols = struct.unpack(">IIII", fh.read(16))
            return np.array(array.array("B", fh.read()), dtype=np.uint8).
↳ reshape(num_data, rows, cols)
```

```

for filename in ['train-images-idx3-ubyte.gz',
                 'train-labels-idx1-ubyte.gz',
                 't10k-images-idx3-ubyte.gz',
                 't10k-labels-idx1-ubyte.gz']:
    download(base_url + filename, filename)

train_images = parse_images('data/train-images-idx3-ubyte.gz')
train_labels = parse_labels('data/train-labels-idx1-ubyte.gz')
test_images = parse_images('data/t10k-images-idx3-ubyte.gz')
test_labels = parse_labels('data/t10k-labels-idx1-ubyte.gz')

return train_images, train_labels, test_images[:1000], test_labels[:1000]

def load_mnist(N_data=None):
    partial_flatten = lambda x: np.reshape(x, (x.shape[0], np.prod(x.shape[1:
→])))
    one_hot = lambda x, k: np.array(x[:, None] == np.arange(k)[None, :],
→dtype=int)
    train_images, train_labels, test_images, test_labels = mnist()
    train_images = (partial_flatten(train_images) / 255.0 > .5).astype(float)
    test_images = (partial_flatten(test_images) / 255.0 > .5).astype(float)
    K_data = 10
    train_labels = one_hot(train_labels, K_data)
    test_labels = one_hot(test_labels, K_data)
    if N_data is not None:
        train_images = train_images[:N_data, :]
        train_labels = train_labels[:N_data, :]

    return train_images, train_labels, test_images, test_labels

def plot_images(images, ax, ims_per_row=5, padding=5, digit_dimensions=(28, 28),
                cmap=matplotlib.cm.binary, vmin=None, vmax=None):
    """Images should be a (N_images x pixels) matrix."""
    N_images = images.shape[0]
    N_rows = np.int32(np.ceil(float(N_images) / ims_per_row))
    pad_value = np.min(images.ravel())
    concat_images = np.full(((digit_dimensions[0] + padding) * N_rows + padding,
                             (digit_dimensions[1] + padding) * ims_per_row +
→padding), pad_value)
    for i in range(N_images):
        cur_image = np.reshape(images[i, :], digit_dimensions)
        row_ix = i // ims_per_row
        col_ix = i % ims_per_row
        row_start = padding + (padding + digit_dimensions[0]) * row_ix

```

```

        col_start = padding + (padding + digit_dimensions[1]) * col_ix
        concat_images[row_start: row_start + digit_dimensions[0],
                        col_start: col_start + digit_dimensions[1]] = cur_image
    cax = ax.matshow(concat_images, cmap=cmap, vmin=vmin, vmax=vmax)
    plt.xticks(np.array([]))
    plt.yticks(np.array([]))
    return cax

```

```

def save_images(images, filename, **kwargs):
    fig = plt.figure(1)
    fig.clf()
    ax = fig.add_subplot(111)
    plot_images(images, ax, **kwargs)
    fig.patch.set_visible(False)
    ax.patch.set_visible(False)
    plt.savefig(filename)

```

```

[ ]: def train_log_regression(images, labels, learning_rate, max_iter):
    """ Used in Q1
        Inputs: train_images, train_labels, learning rate,
        and max num of iterations in gradient descent
        Returns the trained weights (w/o intercept)"""
    N_data, D_data = images.shape
    K_data = labels.shape[1]
    weights = np.zeros((D_data, K_data))

    for i in range(max_iter):
        lny = log_softmax(images, weights)
        E = cross_ent(lny, labels)
        y = np.exp(lny)
        err = y - labels
        a = np.expand_dims(err, axis=1) * np.expand_dims(images, axis=2)
        dE = np.sum(a, axis = 0)
        weights = weights - learning_rate * dE
        #acc = accuracy(lny, labels)
        #print(acc,E)

    w0 = None # No intercept for log-reg
    return weights, w0

```

```

[ ]: def train_gda(images, labels):
    """ Used in Q2
        Inputs: train_images, train_labels
        Returns the trained weights, the intercept, and D x K class means,

```

```

        and D x D common covariance matrix."""
N_data, D_data = images.shape
K_data = labels.shape[1]

# mu
Mu = []
for i in range(K_data):
    tk = labels[:,i]
    muk = np.sum(images.T * tk, axis = 1)/np.sum(tk)
    Mu.append(muk)
Mu = np.array(Mu).T
#print(Mu.shape)

# sigma
Nk = np.sum(labels, axis = 0)
N = np.sum(labels)

Sigma = []
for i in range(K_data):
    nk = Nk[i]
    muk = Mu[:,i]
    Ck = (labels[:,i] == np.ones(N_data))
    xk = images[Ck,:]
    xn_muk = xk-muk
    sigmak = np.cov(xn_muk.T)
    Sigma.append(nk/N*sigmak)
Sigma = sum(Sigma)

# weights

try:
    np.linalg.inv(Sigma)
except:
    Sigma = Sigma + 1/N_data*np.identity(Sigma.shape[0])

weights = np.linalg.inv(Sigma) @ Mu

# w0
pCk = np.mean(labels, axis=0)
w0 = -1/2*np.diagonal(Mu.T @ weights) + np.log(pCk)
#print(w0.shape)

return weights, w0, Mu, Sigma

```

```

[ ]: def log_softmax(images, weights, w0=None):
    """ Used in Q1 and Q2

```

```

        Inputs: images, and weights
        Returns the log_softmax values."""
    if w0 is None: w0 = np.zeros(weights.shape[1])

    N_data, D_data = images.shape

    a = images @ weights + w0
    b = logsumexp(a, axis=1, keepdims=True)
    lny = a - b

    return lny

def cross_ent(log_Y, train_labels):
    """ Used in Q1
        Inputs: log of softmax values and training labels
        Returns the cross entropy."""

    ce = - np.sum(train_labels * log_Y)

    return ce

def predict(log_softmax):
    """ Used in Q1 and Q2
        Inputs: matrix of log softmax values
        Returns the predictions"""

    num = (log_softmax == np.max(log_softmax, axis = 1, keepdims = True))
    pred = np.int32(num)

    return pred

def accuracy(log_softmax, labels):
    """ Used in Q1 and Q2
        Inputs: matrix of log softmax values and 1-of-K labels
        Returns the accuracy based on predictions from log likelihood values"""

    pred = predict(log_softmax)

    ac = []
    for i in range(pred.shape[0]):
        acn = np.sum((pred[i,:]==1) == (labels[i,:]==1)) == pred.shape[1]
        ac.append(acn)

    acc = np.mean(ac)

```

```
return acc
```

```
[ ]: def main():
    N_data = 10000 # Num of samples to be used in training
    # Set this to a small number while experimenting.
    # For log reg, finally use the entire training dataset for training
    ↪(N_data=None).
    # For gda, use as many training samples as your computer can handle.

    train_images, train_labels, test_images, test_labels = load_mnist(N_data)

    # Q1: train logistic regression
    learning_rate, max_iter = .00001, 100
    weights, w0 = train_log_regression(train_images, train_labels,
    ↪learning_rate, max_iter)
    save_images(weights.T, 'weights.png')

    # Q2: train gaussian discriminant
    #weights, w0, Mu, Sigma = train_gda(train_images, train_labels)
    #save_images(Mu.T, 'means.png')

    #new_digit = 0
    #new_images = np.random.multivariate_normal(Mu[:, new_digit], Sigma, 10)
    #save_images((new_images > .5).astype(float), '0_images.png')

    #new_digit = 3
    #new_images = np.random.multivariate_normal(Mu[:, new_digit], Sigma, 10)
    #save_images((new_images > .5).astype(float), '3_images.png')

    # evaluation
    log_softmax_train = log_softmax(train_images, weights, w0)
    log_softmax_test = log_softmax(test_images, weights, w0)

    train_accuracy = accuracy(log_softmax_train, train_labels)
    test_accuracy = accuracy(log_softmax_test[:1000], test_labels[:1000])

    print("Training accuracy is ", train_accuracy)
    print("Test accuracy is ", test_accuracy)

if __name__ == '__main__':
    main()
```