

# hw4-Q1

April 5, 2021

## 1 1. Unsupervised Learning

```
[ ]: %matplotlib inline
import scipy
import numpy as np
import itertools
import matplotlib.pyplot as plt
```

### 1.1 1. Generating the data

First, we will generate some data for this problem. Set the number of points  $N = 400$ , their dimension  $D = 2$ , and the number of clusters  $K = 2$ , and generate data from the distribution  $p(x|z = k) = \mathcal{N}(\mu_k, \Sigma_k)$ . Sample 200 data points for  $k = 1$  and 200 for  $k = 2$ , with

$$\mu_1 = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}, \mu_2 = \begin{bmatrix} 6.0 \\ 0.1 \end{bmatrix} \quad \text{and} \quad \Sigma_1 = \Sigma_2 = \begin{bmatrix} 10 & 7 \\ 7 & 10 \end{bmatrix}$$

Here,  $N = 400$ . Since you generated the data, you already know which sample comes from which class. Run the cell in the IPython notebook to generate the data.

```
[ ]: # TODO: Run this cell to generate the data
num_samples = 400
cov = np.array([[1., .7], [.7, 1.]]) * 10
mean_1 = [.1, .1]
mean_2 = [6., .1]

x_class1 = np.random.multivariate_normal(mean_1, cov, num_samples // 2)
x_class2 = np.random.multivariate_normal(mean_2, cov, num_samples // 2)
xy_class1 = np.column_stack((x_class1, np.zeros(num_samples // 2)))
xy_class2 = np.column_stack((x_class2, np.ones(num_samples // 2)))
data_full = np.row_stack([xy_class1, xy_class2])
np.random.shuffle(data_full)
data = data_full[:, :2]
labels = data_full[:, 2]
```

Make a scatter plot of the data points showing the true cluster assignment of each point using different color codes and shape (x for first class and circles for second class):

```
[ ]: # TODO: Make a scatterplot for the data points showing the true cluster
      ↪ assignments of each point
      # plt.plot(...) # first class, x shape
      # plt.plot(...) # second class, circle shape
      c0 = data[labels==0, :]
      c1 = data[labels==1, :]

      plt.figure()
      plt.plot(c0[:, 0], c0[:, 1], 'x')
      plt.plot(c1[:, 0], c1[:, 1], 'o')
      plt.show()
```

## 1.2 2. Implement and Run K-Means algorithm

Now, we assume that the true class labels are not known. Implement the k-means algorithm for this problem. Write two functions: `km_assignment_step`, and `km_refitting_step` as given in the lecture (Here, `km_` means k-means). Identify the correct arguments, and the order to run them. Initialize the algorithm with

$$\hat{\mu}_1 = \begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \hat{\mu}_2 = \begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}$$

and run it until convergence. Show the resulting cluster assignments on a scatter plot either using different color codes or shape or both. Also plot the cost vs. the number of iterations. Report your misclassification error.

```
[ ]: def cost(data, R, Mu):
      N, D = data.shape
      K = Mu.shape[1]
      J = 0
      for k in range(K):
          J += np.dot(np.linalg.norm(data - np.array([Mu[:, k], ] * N),
      ↪ axis=1)**2, R[:, k])
      return J
```

```
[ ]: # TODO: K-Means Assignment Step
def km_assignment_step(data, Mu):
    """ Compute K-Means assignment step

    Args:
        data: a NxD matrix for the data points
        Mu: a DxK matrix for the cluster means locations

    Returns:
        R_new: a NxK matrix of responsibilities
    """

    # Fill this in:
    N, D = data.shape
```

```

K = Mu.shape[1]
r = np.zeros((N, K))
for k in range(K):
    err = data - np.expand_dims(Mu[:, k], axis=0)
    r[:, k] = np.sum(err**2, axis=1)
arg_min = np.argmin(r, axis=1)
R_new = np.zeros((N, K))
R_new[list(range(N)), arg_min] = 1
return R_new

```

```

[ ]: # TODO: K-means Refitting Step
def km_refitting_step(data, R, Mu):
    """ Compute K-Means refitting step.

    Args:
        data: a NxD matrix for the data points
        R: a NxK matrix of responsibilities
        Mu: a DxK matrix for the cluster means locations

    Returns:
        Mu_new: a DxK matrix for the new cluster means locations
    """
    N, D = data.shape
    K = Mu.shape[1]
    Mu_new = (data.T @ R) / np.sum(R, axis=0)
    return Mu_new

```

```

[ ]: # TODO: Run this cell to call the K-means algorithm
N, D = data.shape
K = 2
max_iter = 100
class_init = np.random.binomial(1., .5, size=N)
R = np.vstack([class_init, 1 - class_init]).T

Mu = np.zeros([D, K])
Mu[:, 1] = 1.
R.T.dot(data), np.sum(R, axis=0)

c = []
for it in range(max_iter):
    R = km_assignment_step(data, Mu)
    Mu = km_refitting_step(data, R, Mu)
    c.append(cost(data, R, Mu))
    #print(it, cost(data, R, Mu))

class_1 = np.where(R[:, 0])
class_2 = np.where(R[:, 1])

```

```
miss_err = np.sum(R[:, 0]==labels)/N
print('The missclassification error is: %f' % miss_err)
```

```
[ ]: # cost vs. the number of iterations
plt.figure()
plt.title('Cost vs. # of iterations')
plt.ylabel('cost')
plt.xlabel('iterations')
plt.plot(np.arange(max_iter), c, 'm-')
plt.show()
```

```
[ ]: # TODO: Make a scatterplot for the data points showing the K-Means cluster
      ↪ assignments of each point
d0 = data[class_1]
d1 = data[class_2]

plt.figure()
plt.plot(d0[:, 0], d0[:, 1], 'x')
plt.plot(d1[:, 0], d1[:, 1], 'o')
plt.show()
```

### 1.3 3. Implement EM algorithm for Gaussian mixtures

Next, implement the EM algorithm for Gaussian mixtures. Write three functions: `log_likelihood`, `gm_e_step`, and `gm_m_step` as given in the lecture. Identify the correct arguments, and the order to run them. Initialize the algorithm with the same initialization as in Q2.1 for the means, and with  $\hat{\Sigma}_1 = \hat{\Sigma}_2 = I$ , and  $\hat{\pi}_1 = \hat{\pi}_2$  for the covariances.

Run the algorithm until convergence and show the resulting cluster assignments on a scatter plot either using different color codes or shape or both. Also plot the log-likelihood vs. the number of iterations. Report your misclassification error.

```
[ ]: def normal_density(x, mu, Sigma):
      return np.exp(-.5 * np.dot(x - mu, np.linalg.solve(Sigma, x - mu))) \
          / np.sqrt(np.linalg.det(2 * np.pi * Sigma))
```

```
[ ]: def log_likelihood(data, Mu, Sigma, Pi):
      """ Compute log likelihood on the data given the Gaussian Mixture
      ↪ Parameters.

      Args:
          data: a NxD matrix for the data points
          Mu: a DxK matrix for the means of the K Gaussian Mixtures
          Sigma: a list of size K with each element being DxK covariance matrix
          Pi: a vector of size K for the mixing coefficients

      Returns:
```

*L: a scalar denoting the log likelihood of the data given the Gaussian Mixture*

```

"""
# Fill this in:
N, D = data.shape
K = Mu.shape[1]
L, T = 0., 0.
for n in range(N):
    T = 0
    for k in range(K):
        T += Pi[k] * normal_density(data[n], Mu[:, k], Sigma[k])
    L += np.log(T)
return L

```

```

[ ]: # TODO: Gaussian Mixture Expectation Step
def gm_e_step(data, Mu, Sigma, Pi):
    """ Gaussian Mixture Expectation Step.

    Args:
        data: a NxD matrix for the data points
        Mu: a DxK matrix for the means of the K Gaussian Mixtures
        Sigma: a list of size K with each element being DxK covariance matrix
        Pi: a vector of size K for the mixing coefficients

    Returns:
        Gamma: a NxK matrix of responsibilities
    """
    # Fill this in:
    N, D = data.shape
    K = Mu.shape[1]
    Gamma = np.zeros((N,K))
    for n in range(N):
        for k in range(K):
            Gamma[n, k] = Pi[k] * normal_density(data[n], Mu[:, k], Sigma[k])
        Gamma[n, :] /= np.sum(Gamma[n, :])
    return Gamma

```

```

[ ]: # TODO: Gaussian Mixture Maximization Step
def gm_m_step(data, Gamma):
    """ Gaussian Mixture Maximization Step.

    Args:
        data: a NxD matrix for the data points
        Gamma: a NxK matrix of responsibilities

    Returns:
        Mu: a DxK matrix for the means of the K Gaussian Mixtures

```

```

        Sigma: a list of size K with each element being DxD covariance matrix
        Pi: a vector of size K for the mixing coefficients
        """
        # Fill this in:
        N, D = data.shape
        K = Gamma.shape[1]
        Nk = np.sum(Gamma, axis=0)
        Mu = np.zeros([D, K])
        Sigma = np.zeros([K, D, D])

        for k in range(K):
            Mu[k] = (data.T @ Gamma)[k] / Nk[k]
            sigmasum = np.zeros([D,D])
            for n in range(N):
                xmmu = data[n:n+1].T - Mu[:,k:k+1]
                sigmasum += Gamma[n,k] * (xmmu @ xmmu.T)
            Sigma[k] = sigmasum / Nk[k]
        Pi = Nk / N
        return Mu, Sigma, Pi

```

```

[ ]: # TODO: Run this cell to call the Gaussian Mixture EM algorithm
N, D = data.shape
K = 2
Mu = np.zeros([D, K])
Mu[:, 1] = 1.
Sigma = [np.eye(2), np.eye(2)]
Pi = np.ones(K) / K
Gamma = np.zeros([N, K]) # Gamma is the matrix of responsibilities

max_iter = 200

loglik = []
for it in range(max_iter):
    Gamma = gm_e_step(data, Mu, Sigma, Pi)
    Mu, Sigma, Pi = gm_m_step(data, Gamma)
    loglik.append(log_likelihood(data, Mu, Sigma, Pi))
    # print(it, log_likelihood(data, Mu, Sigma, Pi)) # This function makes the
    ↪ computation longer, but good for debugging

class_1 = np.where(Gamma[:, 0] >= .5)
class_2 = np.where(Gamma[:, 1] >= .5)

miss_err = np.sum(np.float32(Gamma[:, 0] >= .5) == labels) / N
print('The missclassification error is: %f' % miss_err)

```

```

[ ]: # cost vs. the number of iterations
plt.figure()

```

```
plt.title('log-likelihood vs. # of iterations')
plt.ylabel('log-likelihood')
plt.xlabel('iterations')
plt.plot(np.arange(max_iter), loglik, 'm-')
plt.show()
```

```
[ ]: # TODO: Make a scatterplot for the data points showing the Gaussian Mixture
      ↪ cluster assignments of each point
d0 = data[class_1]
d1 = data[class_2]

plt.figure()
plt.plot(d0[:, 0], d0[:, 1], 'x')
plt.plot(d1[:, 0], d1[:, 1], 'o')
plt.show()
```

#### 1.4 4. Comment on findings + additional experiments

Comment on the results:

- Compare the performance of k-Means and EM based on the resulting cluster assignments.
- Compare the performance of k-Means and EM based on their convergence rate. What is the bottleneck for which method?
- Experiment with 5 different data realizations (generate new data), run your algorithms, and summarize your findings. Does the algorithm performance depend on different realizations of data?

**TODO: Your written answer here**

- EM algorithm gives a much more accurate classification. The missclassification error is lower and the scatterplot is more similar to the true one.
- K-means converge faster than EM algorithm. The bottleneck of the two algorithms are about 5 for k-means and about 25 for EM algorithm.
- The misclassification error of k-means are always higher than that of EM algorithm. So the performance does not depend on different realization of data.

```
[ ]:
```